

HDFS Fundamentals

Hadoop Distributed File System

Professor Anis Koubaa

SE 446
Alfaisal University

Spring 2026

Outline

- 1 HDFS Overview
- 2 HDFS Architecture
- 3 HDFS Read/Write Operations
- 4 Blocks and Replication
- 5 File Formats
- 6 HDFS Commands
- 7 Summary

What is HDFS?

Definition

HDFS (Hadoop Distributed File System) is a distributed storage system designed to store very large files across multiple machines.

Key Features:

- **Distributed**: Files split across many machines
- **Fault-tolerant**: Data replicated for reliability
- **Scalable**: Add more machines as needed
- **Cost-effective**: Uses commodity hardware

HDFS Design Philosophy


Optimized For

- Very large files (GB to PB)
- Streaming data access
- Write once, read many
- Commodity hardware

NOT Optimized For

- Low-latency access
- Many small files
- Random writes
- Interactive queries

HDFS Architecture: Master-Slave



NameNode: The Master

Responsibilities

- Manages file system **namespace** (directory tree)
- Stores **metadata**: file names, permissions, block locations
- Handles **client requests** for file operations
- Tracks **health** of DataNodes via heartbeats

Critical Component!

If NameNode fails, **the entire cluster is unavailable.**

Solutions: Secondary NameNode, HA (High Availability) mode

NameNode Variants: Understanding the Differences

Component	Purpose	When Used
Secondary NameNode	<ul style="list-style-type: none">Periodically merges fsimage + editsCreates checkpoint metadataNOT a backup!	Single NameNode deployments
Standby NameNode	<ul style="list-style-type: none">Hot standby for failoverSynchronized metadataAutomatic takeover on failure	High Availability (HA) mode

Production Recommendation

DataNode: The Workers

Responsibilities

- Store actual **data blocks** on local disk
- Serve **read/write requests** from clients
- Perform **block replication** as instructed by NameNode
- Send **heartbeats** to NameNode every 3 seconds
- Send **block reports** every hour

Commodity Hardware

DataNodes are designed to run on **cheap, commodity servers**.

Failures are expected and handled automatically.

Monitoring DataNodes: Heartbeats & Block Reports

Mechanism	Frequency	Purpose
Heartbeat	Every 3 sec	<ul style="list-style-type: none">• Confirms DataNode is alive• Reports available storage• Receives commands from NameNode
Block Report	Every 1 hour	<ul style="list-style-type: none">• Lists all blocks on this DataNode• Enables metadata verification• Triggers cluster balancing

HDFS Write Path: How Data Gets Stored


- ① **Client requests** to write file to NameNode
- ② **NameNode checks** permissions, creates metadata
- ③ **NameNode selects** DataNodes for block replicas
- ④ **Client streams data** directly to first DataNode
- ⑤ **Replication pipeline:** DataNode-1 → DataNode-2 → DataNode-3
- ⑥ **Acknowledgment** flows back to client
- ⑦ **NameNode updates** metadata

Key Insight

Data never flows through NameNode!

Client writes directly to DataNodes for scalability.

HDFS Write Pipeline Visualization



Pipeline replication: Each DataNode forwards to next while writing locally

HDFS Read Path: How Data Gets Retrieved

- ① **Client requests** file from NameNode
- ② **NameNode returns** block locations (DataNode addresses)
- ③ **Client selects** nearest DataNode for each block
- ④ **Client streams** data directly from DataNodes
- ⑤ **Parallel reads** for different blocks
- ⑥ **Checksum verification** on each block chunk
- ⑦ **Automatic failover** to replica if checksum fails


Performance Benefit

Locality awareness: Client reads from closest DataNode.
Maximizes network bandwidth utilization.

Data Blocks

What is a Block?

Files in HDFS are split into fixed-size **blocks** (default: 128 MB).




Why 128 MB? Minimize seek time, maximize throughput for large files.

Replication Factor

What is Replication?

Each block is copied to **multiple DataNodes** for fault tolerance.

Default replication factor: **3**



Block A is stored on 3 different nodes


Trade-off: More replicas = more fault tolerance, but $3 \times$ storage cost.

Rack Awareness: Smart Replica Placement

Why Rack Awareness?

In real data centers, servers are organized in **racks**.

Entire rack can fail (power, network switch failure).



Default Strategy (3 replicas):

Replica 1: Local rack — Replica 2: Remote rack — Replica 3: Same remote rack

Benefit

Survives entire rack failure while minimizing cross-rack traffic.

Storage Calculation

Example

File size: 500 MB

Replication factor: 3

Total storage used: $500 \times 3 = 1,500 \text{ MB} = 1.5 \text{ GB}$

Formula

$$\text{Total Storage} = \text{File Size} \times \text{Replication Factor}$$

Question: If you store 10 TB of data with replication factor 3, how much disk space do you need?

Data Integrity: Checksums

How HDFS Ensures Data Correctness

HDFS computes **CRC32 checksums** for every 512-byte chunk of data.

Checksum Verification Process:

- ① **On Write:** Compute checksum, store with data
- ② **On Read:** Recompute checksum, compare with stored value
- ③ **If Mismatch:**
 - Report corruption to NameNode
 - Automatically read from another replica
 - NameNode schedules re-replication from good copy

Why It Matters

Protects against:

- Disk corruption (bit rot)
- Network transmission errors

File Formats Comparison

Format	Type	Compression	Best For
CSV	Row-based	None	Simple exchange
JSON	Row-based	None	APIs, configs
Parquet	Column-based	Yes	Analytics
Avro	Row-based	Yes	Streaming
ORC	Column-based	Yes	Hive

For Big Data Analytics

Use **Parquet** — columnar storage with excellent compression.

Why Columnar Storage?

Row-Based (CSV)

ID	Name	Age	Salary
1	Alice	30	75000
2	Bob	25	65000

Column-Based (Parquet)

ID	Name	Age	Salary
1	Alice	30	75000
2	Bob	25	65000

*Query: $\text{SELECT AVG}(\text{Salary})$
Only reads Salary column!*

Parquet Benefits

① Read only needed columns

Faster queries, less I/O

② Better compression

Similar values in a column compress well (10x smaller)

③ Schema embedded

No need for external schema files

④ Predicate pushdown

Filter data at storage level, before loading

Common HDFS Commands

```
# List files in a directory
hdfs dfs -ls /user/data/

# Create a directory
hdfs dfs -mkdir /user/mydata/

# Upload a file from local to HDFS
hdfs dfs -put local_file.csv /user/mydata/

# Download a file from HDFS to local
hdfs dfs -get /user/mydata/file.csv ./local/

# View file contents
hdfs dfs -cat /user/mydata/file.txt

# Delete a file
hdfs dfs -rm /user/mydata/old_file.csv

# Check disk usage
```

Summary: Key Takeaways

- ① **HDFS** = Distributed file system for Big Data
- ② **NameNode** (master) stores metadata; **DataNode** (workers) store data
 - Secondary NameNode: Checkpoints — Standby: HA failover
 - Heartbeats (3s) + Block Reports (1h) for monitoring
- ③ **Write/Read paths**: Client streams directly to/from DataNodes
- ④ Files split into **blocks** (128 MB); **Replication factor 3**
 - Rack awareness for fault tolerance
 - Checksums ensure data integrity
- ⑤ Use **Parquet** for Big Data analytics (columnar, compressed)

Next Week

MapReduce: Distributed data processing paradigm

ExamGPT Quiz Topics

Be prepared to answer questions about:

- Role of NameNode vs DataNode
- Secondary vs Standby NameNode
- Heartbeats vs Block Reports
- HDFS Read/Write paths
- Rack awareness and replica placement
- Storage calculations with replication
- Data integrity (checksums)
- Block size and why it matters
- CSV vs Parquet comparison
- Basic HDFS commands

Quiz in 15 minutes!

Open ExamGPT and complete the Week 2B quiz.

Questions?

Prof. Anis Koubaa
akoubaa@alfaisal.edu