

SE446: Big Data Engineering

Week 3A: MapReduce Concepts

Prof. Anis Koubaa

Alfaisal University - College of Engineering

Spring 2026

Today's Agenda

Recap: Why HDFS Alone Is Not Enough

HDFS provides:

- ✓ Distributed storage
- ✓ Fault tolerance
- ✓ Scalability

But how do we process data?

- Read 10 TB from HDFS?
- Analyze on one machine?
- ✗ Bottleneck!

images/hdfs_to_mapreduce.png


The Big Idea: Move Computation to Data

Traditional Approach


Move data to computation → Network bottleneck

MapReduce Approach

Move computation to data → Process locally, combine results



Traditional



MapReduce

MapReduce: Origin Story

- **2004:** Google publishes landmark paper
- **Problem:** Index the entire web (billions of pages)
- **Solution:** Simple programming model for distributed processing
- **2006:** Doug Cutting implements open-source version → Hadoop

Key Insight

Many data processing tasks follow the same pattern:

- ① Process each record independently (Map)
- ② Group by some key
- ③ Aggregate groups (Reduce)

map(key, value) → list(key', value')

reduce(key', list(value')) → list(value'')

Map Phase:

- Process each input record
- Emit (key, value) pairs
- Runs in parallel across nodes

Reduce Phase:

- Receive all values for a key
- Aggregate/combine values
- Produce final output

Classic Example: Word Count

Input: 3 documents with text

Goal: Count occurrences of each word

Mapper (Python):

```
def mapper(doc_id, text):
    for word in text.split():
        emit(word, 1)
```

Input: "Hello World Hello"

Output:

- (Hello, 1)
- (World, 1)
- (Hello, 1)


Reducer (Python):

```
def reducer(word, counts):
    emit(word, sum(counts))
```

Input: (Hello, [1, 1])

Output: (Hello, 2)

Data Flow: Step by Step



Real-World Example: Chicago Crime Data

Dataset: Chicago crime reports (columns):

- ID, Case Number, Date
- **Primary Type** (THEFT, ASSAULT, BATTERY, ...)
- Description, Location Description
- Arrest (True/False)
- District, Ward, Community Area

Question: How many crimes of each type occurred?

Crime Count: Mapper

```
def crime_mapper(record):
    """
    Input: One crime record (dictionary)
    Output: (crime_type, 1)
    """
    crime_type = record['Primary Type']
    return (crime_type, 1)

# Example
record = {'Primary Type': 'THEFT', 'District': 11, ...}
crime_mapper(record) # Output: ('THEFT', 1)
```

Key Point: Mapper processes ONE record at a time, emits (key, value)

Crime Count: Reducer

```
def crime_reducer(crime_type, counts):
    """
    Input: crime_type (key), list of counts (values)
    Output: (crime_type, total_count)
    """
    total = sum(counts)
    return (crime_type, total)

# Example
crime_reducer('THEFT', [1, 1, 1, 1, 1])
# Output: ('THEFT', 5)
```

Key Point: Reducer receives ALL values for ONE key

Complete Data Flow

Phase	Input	Operation	Output
Input	CSV file	Split by lines	Records
Map	{ "Type": "THEFT" }	Extract type, emit 1	("THEFT", 1)
Map	{ "Type": "ASSAULT" }	Extract type, emit 1	("ASSAULT", 1)
Map	{ "Type": "THEFT" }	Extract type, emit 1	("THEFT", 1)
Shuffle	All map outputs	Group by key	"THEFT" → [1, 1] "ASSAULT" → [1]
Reduce	"THEFT", [1, 1]	Sum	("THEFT", 2)
Reduce	"ASSAULT", [1]	Sum	("ASSAULT", 1)

The Shuffle Phase

Often Overlooked but Critical!

Shuffle is the **most expensive** phase - involves network transfer

What happens during Shuffle:

- ① Mapper outputs are **partitioned** by key
- ② Data is **sorted** by key
- ③ Data is **transferred** to appropriate reducer
- ④ Values for same key are **grouped** together

Example:

- Mapper 1 emits: (A, 1), (B, 2), (A, 3)
- Mapper 2 emits: (B, 4), (A, 5)
- After shuffle: A → [1, 3, 5], B → [2, 4]

Combiners: Local Aggregation

Problem: Shuffle transfers lots of data

Solution: Combine locally before shuffle!

Without Combiner

Mapper 1

(A,1)(A,1)(A,1)



Send 3 values

With Combiner

Mapper 1

(A,1)(A,1)(A,1)

Combiner

(A,3)



Send 1 value!

Note: Combiner must be associative and commutative (like sum, max, min)

Partitioner: Key Distribution

Default: Hash partitioning ($\text{hash}(\text{key}) \% \text{num_reducers}$)

Why customize?

- Ensure related keys go to same reducer
- Balance load across reducers

Example: Analyzing crimes by year

- Key: “2023-THEFT”, “2023-ASSAULT”, “2024-THEFT”
- Partitioner: Extract year, send same year to same reducer
- Result: Each reducer handles one year’s data

MapReduce: Good vs Bad Use Cases

✓ Good for:

- Batch processing
- Large datasets (TB+)
- Embarrassingly parallel tasks
- Aggregations, counts, sums
- Log analysis
- ETL pipelines

✗ Bad for:

- Real-time queries
- Iterative algorithms
- Small datasets
- Interactive analysis
- Graph processing
- Machine learning

Rule of Thumb

If you need results in \downarrow 1 minute, MapReduce is probably wrong choice

MapReduce Limitations

- ① **High Latency:** Minutes to hours for jobs
- ② **Disk I/O:** Intermediate results written to disk
- ③ **Only Two Phases:** Complex workflows need chaining
- ④ **No Iteration:** Each job reads from scratch

The Evolution

MapReduce limitations led to:

- **Spark:** In-memory processing, DAG execution
- **Flink:** True streaming
- **Hive:** SQL on MapReduce

We'll cover these in later weeks!

Key Takeaways

- ① **MapReduce** = programming model for distributed processing
- ② **Map**: Process records in parallel → emit (key, value)
- ③ **Shuffle**: Group by key (automatic, expensive)
- ④ **Reduce**: Aggregate values per key → final result
- ⑤ **Data locality**: Computation moves to data, not vice versa
- ⑥ **Combiner**: Local aggregation to reduce shuffle

Think in Key-Value Pairs!

Every MapReduce problem starts with:

"What should my key be? What should my value be?"

Week 3B: Implementing MapReduce in Python

- Implement mappers and reducers
- Complete crime analysis exercises
- Start Milestone 2

Pre-class preparation:

- Watch: Python MapReduce Tutorial
- Clone team repo and pull latest
- Review Chicago crime dataset