

SE446: Big Data Engineering

Week 3A: MapReduce Concepts

Prof. Anis Koubaa

SE 446
Alfaisal University

https://github.com/aniskoubaa/big_data_course

Spring 2026



جامعة الفيصل

Today's Agenda

- 1 Introduction to MapReduce
- 2 The MapReduce Programming Model
- 3 Detailed Example: Crime Analysis
- 4 Key Concepts
- 5 When to Use MapReduce
- 6 Summary


Recap: Why HDFS Alone Is Not Enough

HDFS provides:

- ✓ Distributed storage
- ✓ Fault tolerance
- ✓ Scalability

But how do we process data?

- Read 10 TB from HDFS?
- Analyze on one machine?
- ✗ Bottleneck!




The Big Idea: Move Computation to Data

Traditional Approach


Move data to computation → Network bottleneck

MapReduce Approach

Move computation to data → Process locally, combine results



Traditional



MapReduce

MapReduce: Origin Story

- **2004:** Google publishes landmark paper
- **Problem:** Index the entire web (billions of pages)
- **Solution:** Simple programming model for distributed processing
- **2006:** Doug Cutting implements open-source version → Hadoop

Key Insight

Many data processing tasks follow the same pattern:

- ① Process each record independently (Map)
- ② Group by some key
- ③ Aggregate groups (Reduce)

MapReduce in One Slide

map(key, value) → list(key', value')

reduce(key', list(value')) → list(value'')

Map Phase:

- Process each input record
- Emit (key, value) pairs
- Runs in parallel across nodes

Reduce Phase:

- Receive all values for a key
- Aggregate/combine values
- Produce final output

Crucial Concept: Mapper Input

Common Confusion

Students often ask: "What is the key passed to the Mapper?"

It is **NOT** the same as the key you emit!

- **Input Format:** Defines how files are read.
- **TextInputFormat (Default):**
 - **Key:** Byte offset of the line (e.g., 0, 104, 256...)
 - **Value:** The actual line of text

Typical Mapper Implementation

We usually **ignore** the input key (offset) and just process the value (text).

The Hidden Step: Shuffle & Sort (System Phase)

Important

Between **Map** and **Reduce**, Hadoop runs a hidden phase called **Shuffle & Sort**.

What Shuffle does (automatically):

- ① Collects mapper outputs (key, value)
- ② Sends each key to the correct reducer
- ③ Sorts keys and groups all values per key

Tiny Example:

Mapper outputs:

(Hello,1) (World,1) (Hello,1)
(World,1) (Hello,1)

After Shuffle:

Hello → [1,1,1]
World → [1,1]

Key Message

You **do not write Shuffle code**.

It is done by the framework (system-level).

Why it matters

Shuffle is often the **slowest** step (network transfer).

Classic Example: Word Count

Input: 3 documents with text

Goal: Count occurrences of each word

Mapper (Python):

```
def mapper(doc_id, text):
    for word in text.split():
        emit(word, 1)
```

Input: "Hello World Hello"

Output:

- (Hello, 1)
- (World, 1)
- (Hello, 1)


Reducer (Python):

```
def reducer(word, counts):
    emit(word, sum(counts))
```

Input: (Hello, [1, 1])

Output: (Hello, 2)

Data Flow: Step by Step



Data Flow Example: Step-by-Step (Word Count)

1) Input Splits (from HDFS)

Split 1: "Hello World"
Split 2: "Hello Hadoop"

2) Map Output (each mapper emits pairs)

Mapper 1 emits:

- (Hello, 1)
- (World, 1)

Mapper 2 emits:

- (Hello, 1)
- (Hadoop, 1)

3) Shuffle & Sort (System-level, automatic)

Group all values with the same key:

- Hello → [1, 1]
- World → [1]
- Hadoop → [1]

4) Reduce Output (final result)

- (Hello, 2)
- (World, 1)
- (Hadoop, 1)

Map creates pairs → Shuffle groups by key → Reduce aggregates

Real-World Example: Chicago Crime Data

Dataset: Chicago Crime Reports

Real-world dataset from the City of Chicago (6GB+, millions of rows).

Key Columns:

- ID, Case Number: Unique identifiers
- Date: Timestamp of incident
- Primary Type: Category (THEFT, BATTERY...)
- Description: Detailed description
- Arrest: Boolean (True/False)

Sample Data (CSV):

Date	Primary Type	Arrest
01/01/2023	THEFT	False
01/01/2023	BATTERY	True
01/02/2023	THEFT	False
01/02/2023	ASSAULT	True
...

The Challenge

How many crimes of each type occurred in total?

Stop & Think: Model the Solution

Before Looking at the Code...

Problem: **Count crimes by type**



What should my KEY be?

Think for 15 seconds...



What should my VALUE be?

Think for 15 seconds...

Stop & Think: Model the Solution

Before Looking at the Code...

Problem: **Count crimes by type**



What should my KEY be?



What should my VALUE be?

→ The Answer

Key = Crime Type (e.g., "THEFT") — **Value** = 1 (for counting)

Crime Count: Mapper

```
def crime_mapper(record):
    """
    Input: One crime record (dictionary)
    Output: (crime_type, 1)
    """
    crime_type = record['Primary Type']
    return (crime_type, 1)

# Example
record = {'Primary Type': 'THEFT', 'District': 11, ...}
crime_mapper(record) # Output: ('THEFT', 1)
```

Key Point: Mapper processes ONE record at a time, emits (key, value)

Crime Count: Reducer

```
def crime_reducer(crime_type, counts):
    """
    Input: crime_type (key), list of counts (values)
    Output: (crime_type, total_count)
    """
    total = sum(counts)
    return (crime_type, total)

# Example
crime_reducer('THEFT', [1, 1, 1, 1, 1])
# Output: ('THEFT', 5)
```

Key Point: Reducer receives ALL values for ONE key

Complete Data Flow (Color-Coded)

Phase	What goes in	What happens	What comes out
Input	CSV file	Split into individual records	One record at a time
Map	{"Primary Type": "THEFT"}	Extract key and emit count	("THEFT", 1)
Shuffle	All mapper outputs <i>(system-controlled)</i>	Group values by key	"THEFT" → [1,1]
Reduce	"THEFT", [1,1]	Aggregate values (sum)	("THEFT", 2)

Data shape evolution: Record → (Key,Value) → Key→[Values] → Final Result

Key Design Cheatsheet

Designing Your Key is 90% of the Battle

Goal	Map Output Key	Map Output Value
Counting (e.g., Word Count)	Item "apple"	1 1
Grouping (e.g., Crimes by Type)	Group ID "THEFT"	Full Record {id: 1, loc: "Street" ...}
Filtering (Map-only job)	- (None)	- (None)
Inverted Index (Search Engine)	Keyword "Hadoop"	Document ID "doc1.txt"

Quick Quiz: Model These Problems (1/2)

For each scenario, identify the Key and Value:

1. Total sales per store

Key: _____

Value: _____

2. Website visits per country

Key: _____

Value: _____

3. Average temperature per city

Key: _____

Value: _____

 Hint: For averages, what EXTRA info do you need?

Quick Quiz: Model These Problems (1/2) - Answers

For each scenario, identify the Key and Value:

1. Total sales per store

Key: `Store ID`

Value: `Sale Amount`

2. Website visits per country

Key: `Country`

Value: `1`

3. Average temperature per city

Key: `City`

Value: `(temp, 1)`

✓ Need `sum` and `count` to compute average!

Quick Quiz: Model These Problems (2/2)

More challenging scenarios:

4. Inverted Index (Search Engine)

Build an index: which documents contain each word?

Key: _____

Value: _____

6. Grouping: All employees per department

Collect all records belonging to the same group.

Key: _____

Value: _____

5. Filtering: Find all orders < \$1000

Only output records matching a condition.

Key: _____

Value: _____

Quick Quiz: Model These Problems (2/2) - Answers

More challenging scenarios:

4. Inverted Index (Search Engine)

Build an index: which documents contain each word?

Key: Word

Value: Document ID

6. Grouping: All employees per department

Collect all records belonging to the same group.

Key: Department ID

Value: Employee record

 Reducer concatenates all employee records!

5. Filtering: Find all orders < \$1000

Only output records matching a condition.

Key: (any / None)

Value: Full record

 Map-only job: emit if condition met, else skip!

Input Splits vs. HDFS Blocks

HDFS Blocks (Storage)

- Physical division of data.
- Fixed size (e.g., 128MB).
- Replicated for safety.

Input Splits (Processing)

- Logical division for Mapper.
- **1 Split = 1 Mapper Instance.**
- Usually, Split Size \approx Block Size.

Exam Key Takeaway

HDFS Blocks are for storage reliability.

Input Splits control the number of mappers.

Visual Example: 400MB File Processing

1. Original File (400MB)

mydata.csv (400MB)



2. HDFS Blocks (Storage)

Block 1
128MB

Block 2
128MB

Block 3
128MB

Block 4
16MB



3. Input Splits (Processing)

Split 1
Mapper 1

Split 2
Mapper 2

Split 3
Mapper 3

Split 4
Mapper 4

Key Points:

- ✓ 4 blocks = 4 replicated chunks
- ✓ 4 splits = 4 parallel mappers
- ✓ Usually 1 block ↔ 1 split

The Shuffle Phase

Often Overlooked but Critical!

Shuffle is the **most expensive** phase - involves network transfer

What happens during Shuffle:

- ① Mapper outputs are **partitioned** by key
- ② Data is **sorted** by key
- ③ Data is **transferred** to appropriate reducer
- ④ Values for same key are **grouped** together

Example:

- Mapper 1 emits: (A, 1), (B, 2), (A, 3)
- Mapper 2 emits: (B, 4), (A, 5)
- After shuffle: A → [1, 3, 5], B → [2, 4]

Reverse Engineering Exercise

After Shuffle, the reducer receives:

- "THEFT" → [1, 1, 1, 1, 1]
- "ASSAULT" → [1, 1, 1]
- "BATTERY" → [1, 1]

Questions:

- ① How many THEFT records were in the input?
- ② What did the Mapper emit for each THEFT record?
- ③ What will the Reducer output for THEFT?

Answers

- ① 5 THEFT records
- ② ("THEFT", 1)
- ③ ("THEFT", 5)

Key Insight

Why Do We Sort Before Reducing?

The "Shuffle" phase guarantees that the Reducer receives keys in **sorted order**.

Why is this mandatory?

- ① **Grouping:** All values for Key K must be contiguous to be passed as a single list (iterator).
- ② **Memory Efficiency:** The Reducer doesn't need to fit all data in RAM. It can "stream" through the sorted keys.

Golden Rule

Reducers *never* see unsorted keys.

Combiners: Local Aggregation

Problem: Shuffle transfers lots of data

Solution: Combine locally before shuffle!

Without Combiner

Mapper 1

(A,1)(A,1)(A,1)



Send 3 values

With Combiner

Mapper 1

(A,1)(A,1)(A,1)

Combiner

(A,3)



Send 1 value!

Note: Combiner must be associative and commutative (like sum, max, min)

Critical Nuance: Combiners are Optional

The "Mini-Reducer" Rule

A Combiner is an **optimization**, NOT a guarantee.

Hadoop may run the combiner:

- **Zero times**: If the spill file is small.
- **Once**: Standard case.
- **Multiple times**: During merge sorts on disk.

Implication for Code

Your Combiner logic must be:

- ① **Associative and Commutative** (Order doesn't matter).
- ② Safe to run repeatedly (Idempotent-ish).
- ③ Input/Output types must match the Reducer input.

Partitioner: Key Distribution

Default: Hash partitioning ($\text{hash}(\text{key}) \% \text{num_reducers}$)

Why customize?

- Ensure related keys go to same reducer
- Balance load across reducers

Example: Analyzing crimes by year

- Key: “2023-THEFT”, “2023-ASSAULT”, “2024-THEFT”
- Partitioner: Extract year, send same year to same reducer
- Result: Each reducer handles one year’s data

The Problem of Data Skew (Hot Keys)

Scenario:

- You count words in Wikipedia.
- Word "the" appears 1 billion times.
- Word "zygote" appears 100 times.

Result: The reducer handling "the" runs for 5 hours.
All other reducers finish in 1 minute. The job waits for the slow reducer ("Straggler").

Impact

1 Reducer bottleneck = Slow Job

Solutions:

- **Salt keys:** "the" → "the-1", "the-2" ...
- **Custom Partitioner**
- **Combine early**

⚡ Design Challenge: Handling Skewed Data

Scenario: Analyze 10 Billion Tweets

- 5 billion mention “Taylor Swift” 🎵
- 5 billion mention everyone else combined

Standard Approach:

- Key = celebrity_name
- Value = 1

✗ **Problem:** One reducer handles 5B values while others are idle!

💡 Your Challenge:

How would you modify the **KEY** to distribute the load?

Hint: Can you “spread” Taylor Swift across multiple reducers?

✓ Solution: Salted Keys

Key = “Taylor Swift-1”, “Taylor Swift-2”, ..., “Taylor Swift-100”

Then run a second MapReduce job to combine the partial counts!

Fault Tolerance: The "Secret Sauce"

Question: What happens if a node crashes during a 10-hour job?

Mapper Failures

- Master detects failure.
- Reschedules the task on another node containing the data replica.
- **✓ Re-execution is safe (Idempotent).**
- *Note: Output is local, so it must be re-run.*

Reducer Failures

- Master detects failure.
- Reschedules the task on another node.
- Retrieves mapper outputs again from source nodes.

User code does NOT need to handle checkpoints or failures. The framework handles it!

One Output File Per Reducer

- HDFS does not support concurrent writes to a single file.
- Each Reducer writes its own part file.
- **Naming Convention:** part-r-00000, part-r-00001, ...

Practical Tip: If you want exactly **one** final output file, set `NumReduceTasks = 1`.
(Warning: This kills parallelism for the Reduce phase!)

MapReduce: Good vs Bad Use Cases

✓ Good for:

- Batch processing
- Large datasets (TB+)
- Embarrassingly parallel tasks
- Aggregations, counts, sums
- Log analysis
- ETL pipelines

✗ Bad for:

- Real-time queries
- Iterative algorithms
- Small datasets
- Interactive analysis
- Graph processing
- Machine learning

Rule of Thumb

If you need results in \downarrow 1 minute, MapReduce is probably wrong choice

MapReduce Limitations

- ① **High Latency:** Minutes to hours for jobs
- ② **Disk I/O:** Intermediate results written to disk
- ③ **Only Two Phases:** Complex workflows need chaining
- ④ **No Iteration:** Each job reads from scratch

The Evolution

MapReduce limitations led to:

- **Spark:** In-memory processing, DAG execution
- **Flink:** True streaming
- **Hive:** SQL on MapReduce

We'll cover these in later weeks!

Key Takeaways

- ① **MapReduce** = programming model for distributed processing
- ② **Map**: Process records in parallel → emit (key, value)
- ③ **Shuffle**: Group by key (automatic, expensive)
- ④ **Reduce**: Aggregate values per key → final result
- ⑤ **Data locality**: Computation moves to data, not vice versa
- ⑥ **Combiner**: Local aggregation to reduce shuffle

🔑 The MapReduce Design Formula

 KEY	What do I want to GROUP BY?
 VALUE	What do I want to AGGREGATE?

✍ Write this on your exam notecard!

Next Session: Hands-On Practice

Week 3B: Implementing MapReduce in Python

- Implement mappers and reducers
- Complete crime analysis exercises
- Start Milestone 2

Pre-class preparation:

- Watch: Python MapReduce Tutorial
- Clone team repo and pull latest
- Review Chicago crime dataset