

SE446: Big Data Engineering

Week 3B: Implementing MapReduce in Python

Prof. Anis Koubaa

SE 446
Alfaisal University

https://github.com/aniskoubaa/big_data_course

Spring 2026



Today's Agenda

- 1 Recap: MapReduce Model
- 2 The MapReduce Emulator
- 3 Python MapReduce Framework
- 4 Word Count Example
- 5 Crime Data Analysis
- 6 Multi-Stage MapReduce
- 7 Common Patterns
- 8 Debugging MapReduce
- 9 Connection to Lab 01
- 10 Summary

Quick Recap: The MapReduce Pipeline



Map Phase

$$(k_1, v_1) \rightarrow [(k_2, v_2)]$$

→ Process each record independently and emit key-value pairs.

Reduce Phase

$$(k_2, [v_2]) \rightarrow [(k_3, v_3)]$$

→ Aggregate all values associated with the same key.

Two Ways to Practice MapReduce

❑ Local Emulator (Today)

- Pure Python script
- Runs entirely on your laptop
- Uses loops instead of distributed nodes
- **Purpose:** Learn the algorithm logic

✓ Fast iteration, easy debugging

☰ Real Hadoop Cluster (Lab 01)

- Hadoop Streaming utility
- Runs on multi-node HDFS cluster
- Parallel processing across machines
- **Purpose:** Learn production deployment

✓ Real-world scalability

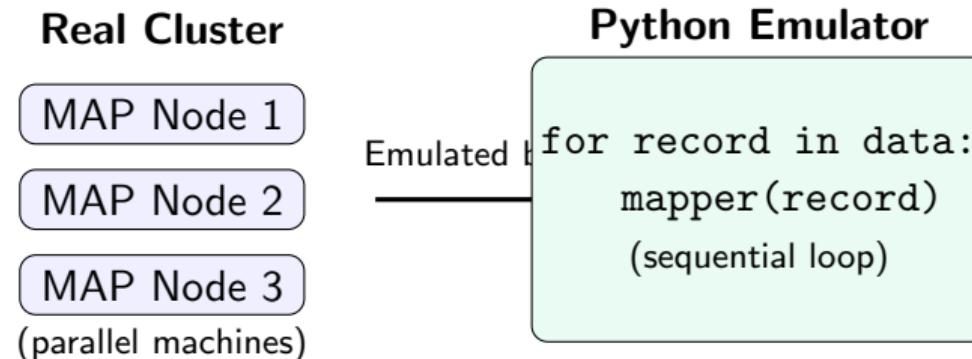
Key Insight

The **mapper** and **reducer** logic you write is **identical** in both approaches. Only the execution environment differs!

Why Use an Emulator?

Problem: Setting up a Hadoop cluster takes time and infrastructure.

Solution: Emulate the MapReduce behavior locally.



Bottom line: Once your algorithm works locally, it will work on a 1000-node cluster!

How the Emulator Works

The Python emulator simulates all three MapReduce phases:

| Phase | Real Cluster | Python Emulator |
|---------|-------------------------------|---------------------------|
| Map | Parallel nodes process chunks | for loop iterates records |
| Shuffle | Network transfers data by key | defaultdict groups by key |
| Reduce | Parallel reducers aggregate | for loop iterates keys |

The Trade-Off

- **Lost:** True parallelism, fault tolerance, distributed storage
- **Gained:** Simplicity, debuggability, instant feedback

Building the Emulator: Overview

We will build a simple Python function called `map_reduce()` that:

① Takes inputs:

- `data`: A list of records (e.g., lines, dictionaries)
- `mapper`: A function you write
- `reducer`: A function you write

② Executes the Map → Shuffle → Reduce pipeline

③ Returns the final aggregated results

Your job as a developer: Write the `mapper` and `reducer` functions.

The framework's job: Handle the shuffle and coordination.

The Emulator Code

```
from collections import defaultdict

def map_reduce(data, mapper, reducer):
    """Python MapReduce Emulator - Simulates distributed processing locally."""

    # ===== MAP PHASE (simulates parallel map nodes) =====
    mapped = []
    for record in data:
        result = mapper(record)
        if result: mapped.append(result)

    # ===== SHUFFLE PHASE (simulates network grouping) =====
    shuffled = defaultdict(list)
    for key, value in mapped:
        shuffled[key].append(value)

    # ===== REDUCE PHASE (simulates parallel reducers) =====
    results = []
    for key, values in shuffled.items():
        results.append(reducer(key, values))

    return results
```

Word Count: The "Hello World" of Big Data

Goal: Count how many times each word appears in a dataset.

Think in Key-Value pairs:

- **Key:** The word itself
- **Value:** The number 1 (one occurrence)

Mapper logic:

For each word, emit (word, 1)

Reducer logic:

Sum all the 1s for each word

Example trace:

- ① Input: "hello world hello"
- ② Map: [("hello",1), ("world",1), ("hello",1)]
- ③ Shuffle: {"hello": [1,1], "world": [1]}
- ④ Reduce: [("hello",2), ("world",1)]

Word Count: The Code

```
# 1. Prepare Data
documents = [
    "Hello World Hello",
    "World of Big Data",
    "Hello Big Data World"
]

# 2. Input Splitting
words = [word for doc in documents
         for word in doc.split()]

# 3. Mapper: Emit (word, 1)
def word_mapper(word):
    return (word.lower(), 1)

# 4. Reducer: Sum values
def count_reducer(word, counts):
    return (word, sum(counts))

# 5. Run MapReduce
results = map_reduce(
    words,
    word_mapper,
    count_reducer
)

print(results)
# [('hello', 3),
#  ('world', 3), ...]
```

Loading Chicago Crime Data

Dataset: Chicago Crime Records

Key Columns:

- Primary Type: THEFT, ASSAULT, BATTERY, etc.
- District: Police district number (1-25)
- Arrest: True/False - was someone arrested?
- Location Description: STREET, RESIDENCE, etc.

Note: We generate 1,000 synthetic records in the notebook to ensure consistency across all students.

Sample Record (as dict):

```
{  
    'ID': 1,  
    'Primary Type': 'THEFT',  
    'District': 3,  
    'Arrest': False,  
    'Location Description':  
        'STREET'  
}
```

Exercise 1: Count Crimes by Type

Goal: How many crimes of each type occurred?

Design Decision:

- **Key:** Crime type (e.g., "THEFT")
- **Value:** 1 (for counting)

Why this works:

Shuffle groups all records with the same crime type. Reducer sums the 1s to get the total count.

Expected Output:

| | |
|-----------------|-----|
| THEFT | 300 |
| BATTERY | 250 |
| ASSAULT | 150 |
| CRIMINAL DAMAGE | 100 |
| BURGLARY | 100 |
| OTHER | 100 |

Exercise 1: The Code

```
# 1. Prepare Data
crime_records = crimes_df.to_dict('records')

# 2. Mapper: Emit (type, 1)
def crime_type_mapper(record):
    return (record['Primary Type'], 1)

# 3. Reducer: Sum counts (from previous example)
def count_reducer(crime_type, counts):
    return (crime_type, sum(counts))

# 4. Run MapReduce
crime_counts = map_reduce(crime_records, crime_type_mapper, count_reducer)

print(crime_counts[:3])
# [('THEFT', 300), ('BATTERY', 250), ...]
```

Exercise 2: Count Crimes per District

Goal: Which police districts have the most crime?

Design Decision:

- **Key:** District number
- **Value:** 1 (for counting)

Observation:

The only change from Exercise 1 is what we use as the key! The reducer stays the same.

This is the power of MapReduce: Change the key, get different insights.

Code Change:

```
# OLD (count by type)
crime_type = record['Primary Type']
return (crime_type, 1)
```

```
# NEW (count by district)
district = record['District']
return (district, 1)
```

Exercise 2: The Code

```
# Mapper: extract district as key
def district_mapper(record):
    district = record['District']
    return (district, 1)

# Same reducer (count pattern)
results = map_reduce(crime_records, district_mapper, count_reducer)

# Sort by count (descending)
sorted_districts = sorted(results, key=lambda x: x[1], reverse=True)

print("Top 5 Districts by Crime Count:")
for district, count in sorted_districts[:5]:
    print(f"  District {district}: {count} crimes")
```

Exercise 3: Filter - Crimes with Arrests

Goal: Count only crimes where an arrest was made.

The Filter Pattern:

- Mapper returns None for unwanted records
- Framework skips None values
- Only matching records reach the reducer

When to use:

Any time you want to analyze a **subset** of your data.

Logic Flow:

```
IF record['Arrest'] == True:  
    EMIT (crime_type, 1)  
ELSE:  
    EMIT None # Skip this record
```

Key Pattern

Return None from mapper to filter records. The framework handles the rest.

Exercise 3: The Code

```
def arrest_mapper(record):
    """
    Only emit crimes where arrest was made
    """
    if record['Arrest'] == True:
        return (record['Primary Type'], 1)
    return None # Filter out - no arrest

# Run MapReduce
arrest_counts = map_reduce(crime_records, arrest_mapper,
                           count_reducer)

print("Crimes with Arrests:")
for crime_type, count in sorted(arrest_counts, key=lambda x: x[1],
                                 reverse=True)[:5]:
    print(f" {crime_type}: {count} arrests")
```

Exercise 4: Aggregation - Arrest Rate by Type

Goal: What percentage of each crime type results in an arrest?

Challenge:

We need two pieces of information:

- ① Total number of crimes
- ② Number of arrests

Solution: Emit a **tuple** as the value!

(crime_type, (arrested, 1))

- arrested: 1 if arrest, else 0
- 1: always 1 (for total count)

Reducer Logic:

```
total_arrests = sum(v[0] for v in values)
total_crimes = sum(v[1] for v in values)
rate = total_arrests / total_crimes * 100
```

Example:

THEFT receives values:

[(1,1), (0,1), (0,1), (1,1)]
→ 2 arrests / 4 total = 50%

Exercise 4: The Code

```
def arrest_stats_mapper(record):
    """Emit (crime_type, (arrested, total))"""
    crime_type = record['Primary Type']
    arrested = 1 if record['Arrest'] else 0
    return (crime_type, (arrested, 1))

def arrest_rate_reducer(crime_type, values):
    """Calculate arrest rate"""
    total_arrests = sum(v[0] for v in values)
    total_crimes = sum(v[1] for v in values)
    rate = (total_arrests / total_crimes) * 100
    return (crime_type, round(rate, 1))

results = map_reduce(crime_records, arrest_stats_mapper, arrest_rate_reducer)

print("Arrest Rates by Crime Type:")
for crime_type, rate in sorted(results, key=lambda x: x[1], reverse=True)[:5]:
    print(f"  {crime_type}: {rate}%")
```

Chaining MapReduce Jobs

Problem: Find the top 5 crime types.

Why is this hard?

- MapReduce processes each key **independently**
- Finding "top 5" requires comparing **all** keys together

Solution: Two-stage MapReduce

① Job 1: Count crimes by type

- Map: (record) \rightarrow (type, 1)
- Reduce: (type, [1,1,1...]) \rightarrow (type, count)

② Job 2: Find top 5

- Map: (type, count) \rightarrow ("all", (type, count)) $// \text{ same key!}$
- Reduce: ("all", [(type, count)...]) \rightarrow sorted top 5

Trick

Using a **constant key** (like "all") sends all data to a single reducer!

Multi-Stage Example

```
# Stage 1: Count by type
crime_counts = map_reduce(crime_records, crime_type_mapper, count_reducer)

# Stage 2: Find top 5
def top_mapper(item):
    """Send all to same reducer with dummy key"""
    crime_type, count = item
    return ("all", (crime_type, count)) # Key="all" sends all to one reducer

def top_reducer(key, values):
    """Sort and take top 5"""
    sorted_values = sorted(values, key=lambda x: x[1], reverse=True)
    return (key, sorted_values[:5])

top_5 = map_reduce(crime_counts, top_mapper, top_reducer)

print("Top 5 Crime Types:")
for crime_type, count in top_5[0][1]:
    print(f" {crime_type}: {count}")
```

MapReduce Design Patterns

| Pattern | Map Output | Reduce Operation |
|----------|--------------------------|------------------|
| Counting | (key, 1) | sum(values) |
| Sum | (key, value) | sum(values) |
| Average | (key, (value, 1)) | sum(v)/sum(c) |
| Max/Min | (key, value) | max/min(values) |
| Filter | (key, value) or None | pass through |
| Distinct | (value, None) | emit key only |
| Top N | (constant, (key, value)) | sort and slice |

Golden Rule: Always ask yourself:

- ➊ **What is my key?** → What do I want to group by?
- ➋ **What is my value?** → What do I want to aggregate?

Common Mistakes

① Forgetting to handle None

- Mapper can return None to filter
- Framework must check for None

② Wrong key choice

- Key determines grouping
- Choose key based on what you want to aggregate

③ Non-hashable keys

- Keys must be hashable (strings, numbers, tuples)
- Lists, dicts cannot be keys

④ Type mismatches

- Ensure values are consistent types
- `sum([1, 1, '1'])` will fail!

Debugging Tips

```
# Add debugging to mapper
def debug_mapper(record):
    result = crime_type_mapper(record)
    print(f"Mapper: {record['ID']} -> {result}")
    return result

# Test with small sample first
sample = crime_records[:5]
test_results = map_reduce(sample, debug_mapper, count_reducer)

# Verify shuffle output
def debug_map_reduce(data, mapper, reducer):
    # ... map phase ...
    print(f"After map: {len(mapped)} pairs")
    # ... shuffle phase ...
    print(f"After shuffle: {len(shuffled)} keys")
    for key in list(shuffled.keys())[:3]:
        print(f"Key: {key}, Value: {shuffled[key]}")
```

From Emulator to Real Cluster (Lab 01)

What changes when moving to Hadoop Streaming?

| Aspect | Python Emulator | Hadoop Streaming |
|-------------|-----------------|------------------|
| Input | Python list | HDFS files |
| Mapper I/O | Function call | stdin → stdout |
| Shuffle | defaultdict | Hadoop framework |
| Reducer I/O | Function call | stdin → stdout |
| Output | Python list | HDFS files |

Key insight: The **logic** stays the same! Only the **I/O mechanism** changes.

In Lab 01, you will:

- Write `mapper.py` that reads from `sys.stdin`
- Write `reducer.py` that reads sorted input from `sys.stdin`
- Submit to the real Hadoop cluster using `mapred streaming`

Today's Lab Tasks

Lab 01: MapReduce with Hadoop Streaming:

- ① Implement word count mapper and reducer
- ② Test locally using Linux pipes: `cat test.txt | python3 mapper.py | sort | python3 reducer.py`
- ③ Run on real Hadoop cluster

Notebook Practice:

- ① Count crimes by type
- ② Find district with most crimes
- ③ Calculate arrest rate by crime type
- ④ (Bonus) Find top 5 crime locations

Deliverables:

- Complete lab notebook
- Commit to team GitHub repo

Summary

What You Learned Today

- **Emulator concept:** Practice MapReduce locally before cluster deployment
- **Framework:** Built a reusable `map_reduce()` function
- **Patterns:** Count, filter, average, multi-stage
- **Debugging:** Test with small samples, add print statements

Key Takeaway:

Think: **What is my key? What is my value?**

Next Week: Milestone 2 work session + MapReduce on Hadoop cluster